

Computational Finance



Plotting Basics

- Plotting in (scientific) Python is mostly done via the `matplotlib` library ([Documentation](#)), which is inspired by the plotting facilities of Matlab®.
- Its main plotting facilities reside in its `pyplot` module. Usually imported as

```
In [2]: import matplotlib.pyplot as plt
        %matplotlib inline
```

- The second line is an [ipython magic](#). It makes plots appear inline in the notebook.
- The `seaborn` library ([Documentation](#)) provides higher-level statistical visualizations:


```
In [3]: import seaborn as sns
```

- Finally, `statsmodels` is useful for QQ plots (see below):

```
In [4]: import statsmodels.api as sm
```


- I will only give a brief introduction to matplotlib here. However, the code for all graphs shown below is included in the notebook (though sometimes hidden in slide mode), and should be studied.
- The fundamental object in matplotlib is a `Figure`, inside of which reside `subplots` (or `axes`).
- To create a new figure, add an axis, and plot to it:

```
In [5]: #with the inline backend, these need to be in the same cell.  
fig=plt.figure(figsize=(6,3)) #create a new empty figure object. Size is optional.  
ax1=fig.add_subplot(121)  #(1x2) axes and make the first one current (what plt.* commands operate on)  
ax2=fig.add_subplot(122)  #(1x2) axes and make the second one current  
ax1.plot(range(10))  
ax2.plot(range(10,0,-1));
```




- By default, matplotlib plots into the current axis, creating one (and a figure) if needed. Using the convenience method `subplot`, this allows us to achieve the same without explicit reference to figures and axes:

```
In [6]: plt.subplot(121)
plt.plot(range(10))
plt.subplot(122)
plt.plot(range(10, 0, -1));
```



- To plot two vectors x and y against each other:

```
In [7]: import numpy as np
x=np.linspace(-10,10,100)
y=x**2
plt.plot(x,y,'r--') #dashed red line; see table on p. 114
plt.xlabel('$x$') #LaTeX equations can be included by enclosing in $$
plt.ylabel('$y$')
plt.title('A Parabola')
plt.legend(['$f(x)$']); #Expects a list of strings
plt.xlim(xmin=-8, xmax=8); #axis limits
#plt.savefig('filename.svg') #to save a plot to disk
```



Risk Measures

Introduction

- The Basel Accords mandate that financial institutions report the risk associated with their positions, so that regulators may check the adequacy of the economic capital as a buffer against market risk.
- Reporting is in the form of a *risk measure*, which condenses the risk of a position into a single number.
- Currently, the mandated measure is *Value at Risk* (VaR), but there are debates of replacing it with an alternative (*Expected Shortfall*).
- Banks are allowed to use their own, internal models for the computation of VaR, but the adequacy of these models should be *backtested*.

Value at Risk

- Consider a portfolio with value V_t and daily (simple) returns R_t .
- Define the one-day loss on the portfolio as

$$Loss_{t+1} = -[V_{t+1} - V_t].$$

- I will distinguish between the dollar Value at Risk (an amount) and the return Value at Risk (a percentage). When unqualified, I mean the latter.
- The one-day $100p\%$ dollar Value at Risk VaR_{t+1}^p is the loss on the portfolio that we are $100(1 - p)\%$ confident will not be exceeded. The Basel committee prescribes $p = 0.01$.

- The *return Value at risk* VaR_{t+1}^p expresses $\$VaR_{t+1}^p$ as a percentage of the portfolio value:

$$VaR_{t+1}^p = \frac{\$VaR_{t+1}^p}{V_t}.$$

- Hence


$$\Pr(R_{t+1} < -VaR_{t+1}^p) = p,$$

because

$$R_{t+1} = -\frac{\$Loss_{t+1}}{V_t}.$$

This holds approximately for log returns, too.



- Thus VaR_{t+1}^p is minus the $100p$ th *percentile* (or minus the p th *quantile*) of the return distribution.



Asset Returns: Stylized Facts

- Stylized facts about asset returns include
 - Lack of autocorrelation
 - Leverage effects
 - Heavy tails of return distribution
 - Volatility clustering
- These need to be taken into account when creating VaR forecasts.

```
In [9]: import pandas as pd
import pandas_datareader.data as web
import scipy.stats as stats #The book likes to import it as `scs`
p=web.DataReader("^GSPC", 'yahoo', start='1/1/2013', end='10/12/2017')['Adj Close']
r=np.log(p)-np.log(p).shift(1)
r.name='Return'
r=r[1:] #remove first observation (NaN)
plt.figure(figsize=(12,4))
plt.subplot(121)
sns.distplot(r, kde=False, fit=stats.norm) #histogram overlaid with fitted normal density
plt.subplot(122)
r.plot() #note that this is a pandas method! looks prettier than plt.plot(r)
plt.savefig('img/stylizedfacts.svg') #save to file
plt.close()
```



VaR Methods: Unconditional

Non-parametric: Historical Simulation


- Historical simulation assumes that the distribution of tomorrow's portfolio returns is well approximated by the empirical distribution (histogram) of the past m observations $\{R_t, R_{t-1}, \dots, R_{t+1-m}\}$.
- This is as if we draw, with replacement, from the last m returns and use this to simulate the next day's return distribution.
- The estimator of VaR is given by minus the p th *sample quantile* of the last m portfolio returns, i.e., $\widehat{VaR}_{t+1}^p = -R_p^m$, where R_p^m is the number such that $100p\%$ of the observations are smaller than it.

- In Python, we can use NumPy's `quantile` method, or the `percentile` function (or `nanpercentile` which ignores NaNs). Hilpisch uses `scoreatpercentile`, but that is deprecated.

```
In [10]: VaR_hist=-r.quantile(.01) #Alternatively, VaR=np.percentile(r,1)  
VaR_hist
```

```
Out[10]: 0.02131716077914799
```

```
In [11]: ax=r.hist(bins=30) #another pandas method: histogram with 30 bins  
ax.set_xticks([-VaR_hist])  
ax.set_xticklabels(['$-VaR_{t+1}^{0.01}$ =  $-%4.3f$ ' %VaR_hist]) #4.3f means 4 digits, of which 3 decimals  
plt.title('Historical VaR')  
plt.savefig('img/var_hist.svg')  
plt.close()
```



- Problem: Last year(s) of data not necessarily representative for the next few days (because of, e.g., volatility clustering).
- Exacerbated by the fact that a large m is required to compute 1% VaR with any degree of precision (only 1% of the data are really used).

Parametric: Normal and t Distributions

- Another simple approach is to assume $R_{t+1} \sim N(\mu, \sigma^2)$, and to estimate μ and σ^2 from historical data (for daily data, $\mu \approx 0$). The VaR is then determined from

$$\begin{aligned}\Pr(R_{t+1} < -VaR_{t+1}^p) &= \Pr\left(\frac{R_{t+1} - \mu}{\sigma} < \frac{-VaR_{t+1}^p - \mu}{\sigma}\right) \\ &= \Pr\left(z_{t+1} < \frac{-VaR_{t+1}^p - \mu}{\sigma}\right) \\ &= \Phi\left(\frac{-VaR_{t+1}^p - \mu}{\sigma}\right) = p,\end{aligned}$$

where $\Phi(z)$ is the cumulative standard normal distribution.

- Thus,


$$VaR_{t+1}^p = -\mu - \sigma\Phi^{-1}(p),$$

where $\Phi^{-1}(p)$ is the inverse distribution function of the standard normal, a.k.a. the *percentage point function* (ppf).

- In Python:


```
In [12]: mu, sig=stats.norm.fit(r) #fit a normal distribution to r
         VaR_norm=-mu-sig*stats.norm.ppf(0.01)
         VaR_norm
```

```
Out[12]: 0.017257996794445292
```



- Problems:
 - Variance of the past year(s) of data not necessarily representative for the future.
 - Returns typically have heavier tails than the normal.
- The solution to the second point is to use another distribution. The Student's t distribution is a popular choice.

- The Student's t distribution with ν degrees of freedom, t_ν , is well known from linear regression as the distribution of t -statistics, where $\nu = T - k$.
- Can be generalized to allow $\nu \in \mathbb{R}_+$.
- Smaller values of ν correspond to heavier tails. As $\nu \rightarrow \infty$, we approach the $N(0, 1)$ distribution.
- It only has moments up to but not including ν :
 - The mean is finite and equal to zero if $\nu > 1$.
 - The variance is finite and equal to $\nu/(\nu - 2)$ if $\nu > 2$.
 - The excess kurtosis is finite and equal to $6/(\nu - 4)$ if $\nu > 4$.
- The distributions are symmetric around 0, hence mean and skewness are 0 if they exist.




- For financial applications, we need to allow for a non-zero mean, and a variance different from $\nu/(\nu - 2)$.
- This is achieved by introducing a *location parameter* m and a *scale parameter* h . We'll write $f_\nu(x; m, h)$ for the resulting density, $F_\nu(x; m, h)$ for the distribution function, and $F_\nu^{-1}(p; m, h)$ for the percentage point function.
- Note that if $x \sim t_\nu(m, h)$, $\nu > 2$, then $\mathbb{E}[x] = m$ and $\text{var}[x] = h^2 \nu/(\nu - 2)$.
- The VaR becomes

$$\text{VaR}_{t+1}^p = -m - hF_\nu^{-1}(p; 0, 0).$$

- In Python:

```
In [15]: df, m, h=stats.t.fit(r) #fit a location-scale t distribution to r
         VaR_t=-m-h*stats.t.ppf(0.01, df)
         VaR_t
```

```
Out[15]: 0.021244629280331669
```



- There are several ways to assess whether a distributional assumption is adequate.
- One is to use a *goodness of fit test*. Many such tests exist.
- Hilpisch discusses the D'Agostino-Pearson test, available as `stats.normaltest`.
- Here we use the Jarque-Bera test. The test statistic is

$$JB = N \left(S^2/6 + (K - 3)/24 \right) ,$$

where S and K are respectively the sample skewness and kurtosis.

- Intuitively, it tests that the skewness and excess kurtosis are zero.
- It is distributed as χ^2_2 under the null of normality. The 5% critical value is

```
In [17]: stats.chi2.ppf(0.95, 2)
```


```
Out[17]: 5.9914645471079799
```

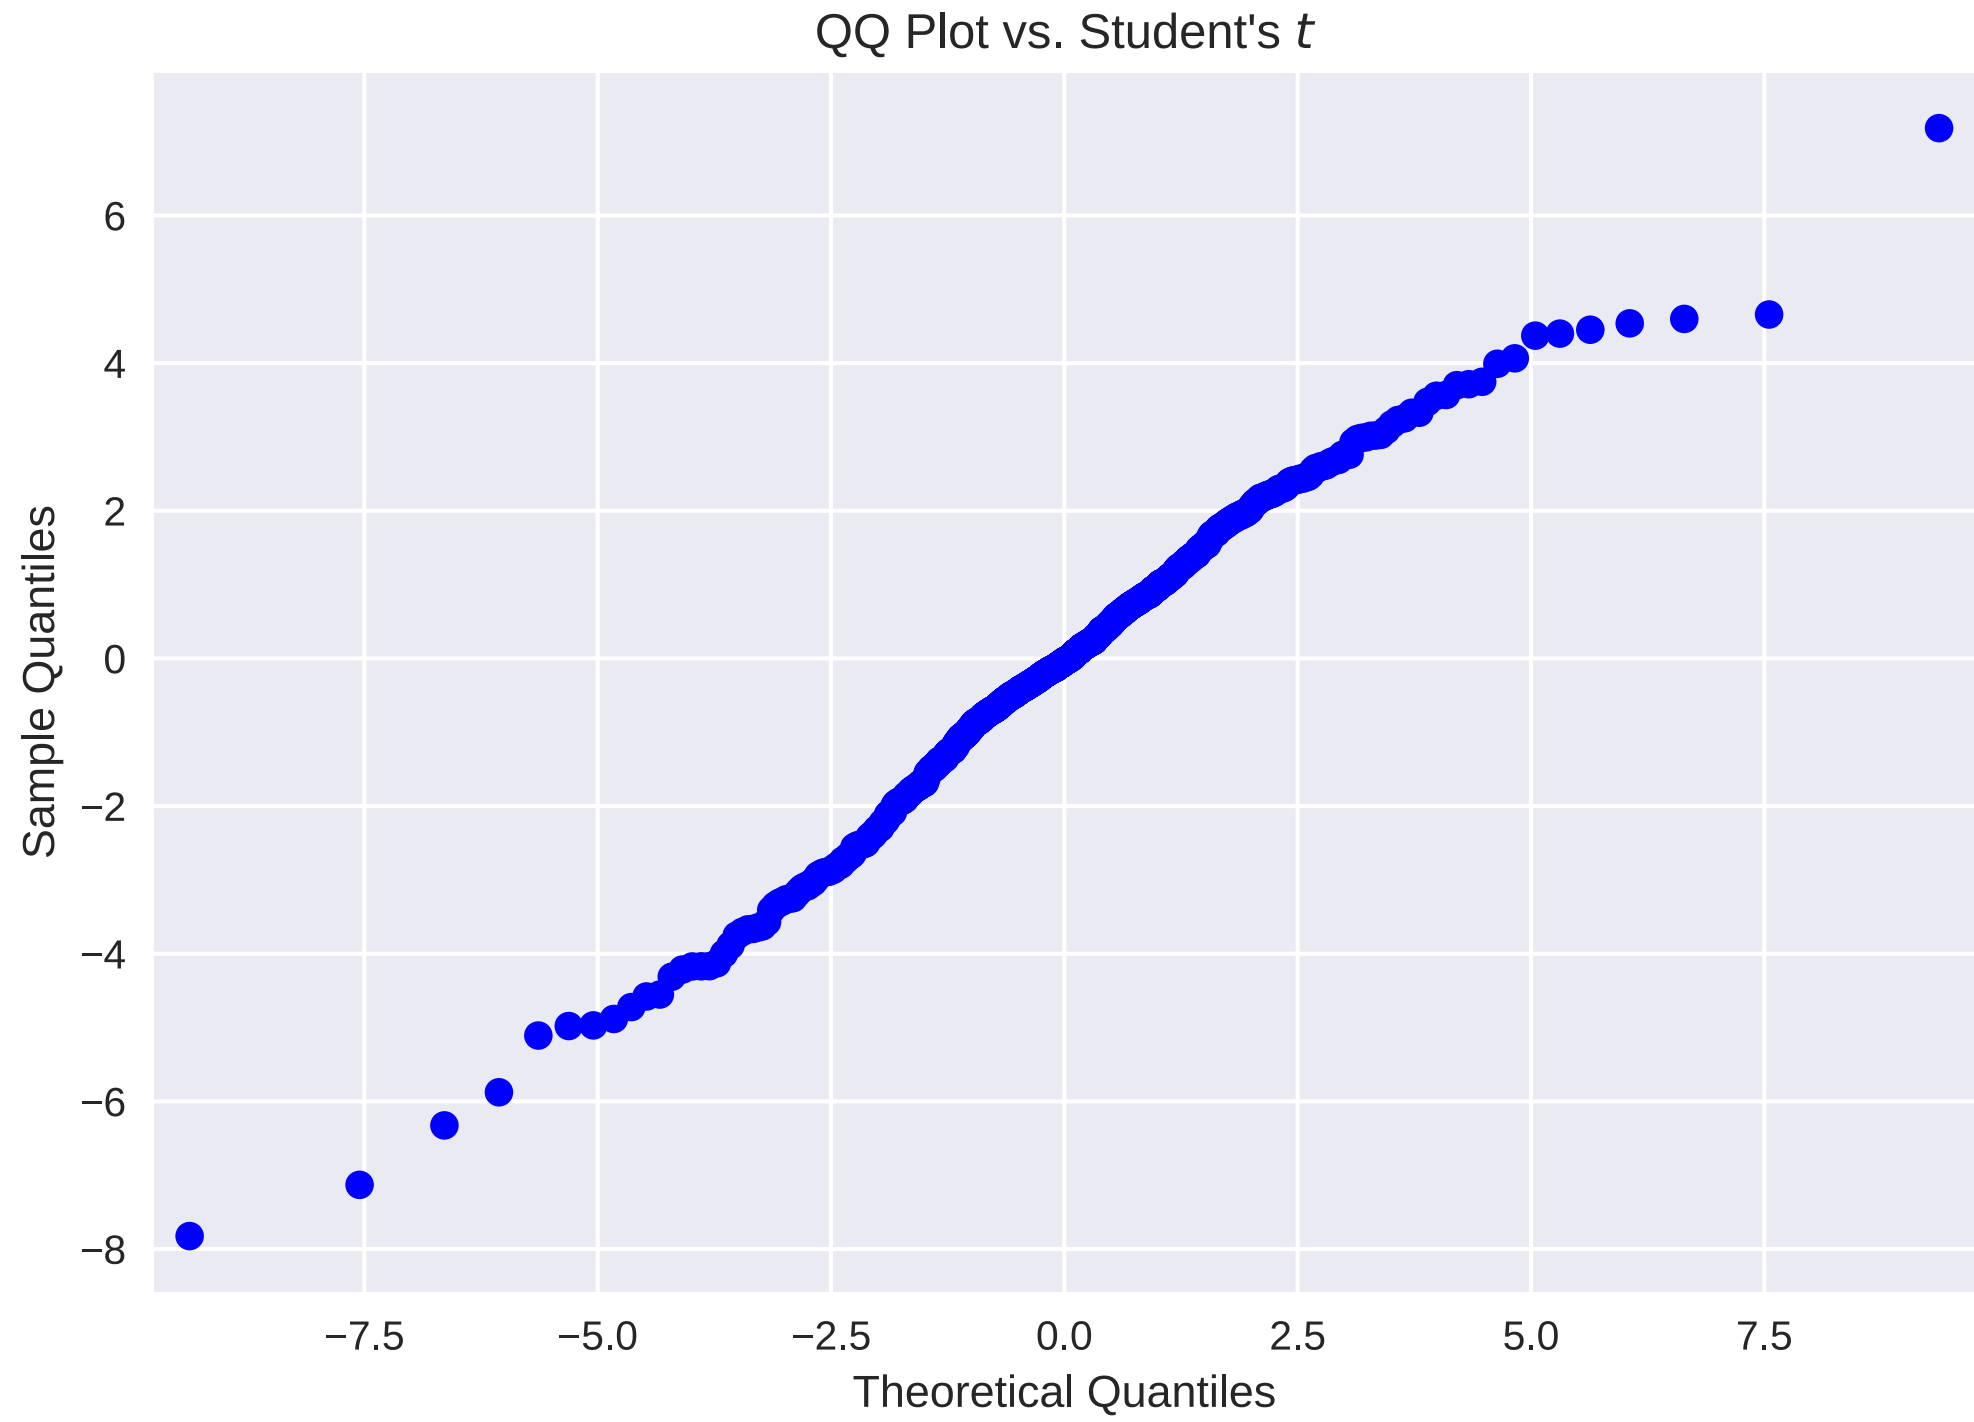
- In Python:

```
In [18]: stats.jarque_bera(r) #returns (JB, p-val)
```

```
Out[18]: (410.78923631633768, 0.0)
```

- Another option is to use a QQ-plot (quantile-quantile plot).
- It plots the empirical quantiles against the quantiles of a hypothesized distribution, e.g. $\Phi^{-1}(p)$ for the normal.
- If the distributional assumption is correct, then the plot should trace out the 45 degree line.





VaR Methods: Filtered

- All methods discussed so far share one drawback: they assume that the volatility is constant, at least in the estimation (and forecast) period.
- Implicitly, the Normal and Student's t method use the *historical volatility*:

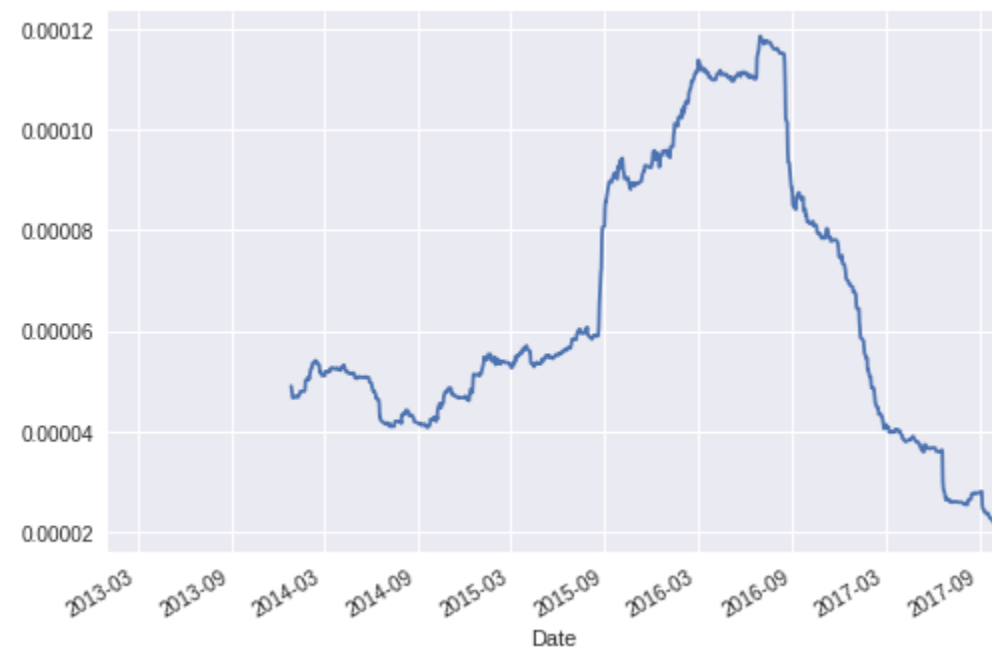
$$\sigma_{t+1,HIST}^2 = \frac{1}{m} \sum_{j=0}^{m-1} R_{t-j}^2.$$

(Note: volatility usually means standard deviation, not variance. I'll be sloppy here).

- Here we assumed a zero mean, which is realistic for daily returns.
- Some adaptability is gained by choosing a smaller m such as 250 (one trading year), but there is a tradeoff because doing so decreases the sample size.
- A general solution requires a *volatility model*, which will be discussed in *Advanced Risk Management*.

- A Pandas Series object has a `rolling` method that can be used to construct historical volatilities for an entire series, using, at each day, the past m observations.
- The method returns a special window object that in turn has a method `var` (for variance).

```
In [21]: sig2_hist=r.rolling(window=250).var()  
sig2_hist.plot();
```




- A partial solution to the drawbacks of historical volatility is given by the RiskMetrics model, which is a special case of a more general framework known as *GARCH* models.
- The idea is to replace the equally weighted moving average used in historical volatility by an exponentially weighted moving average (EWMA):

$$\begin{aligned}\sigma_{t+1,EWMA}^2 &= (1 - \lambda) \sum_{j=0}^{\infty} \lambda^j R_{t-j}^2 \\ &= \lambda \sigma_{t,EWMA}^2 + (1 - \lambda) R_t^2, \quad 0 < \lambda < 1.\end{aligned}$$

- This means that observations further in the past get a smaller weight.
- Smaller λ means faster downweighting; for $\lambda \rightarrow 1$ we approach historical volatility (with an expanding window). For daily data, RiskMetrics recommends $\lambda = 0.94$.
- In practice we do not have $R_{t-\infty}$, but the second equation can be started up by an initial estimate / guess $\sigma_{0,EWMA}^2$.

- The ewm (exponentially moving average) method of a Pandas Series can be used to achieve something similar (the exact definition is slightly different, see [here](#)).
- As before, the method returns a window object that has a var method.

```
In [22]: sig2_ewma=r.ewm(alpha=0.06).var() #alpha=(1-lambda)
sig2_ewma.plot();
```



- The idea behind a filtered VaR method is to decompose the returns as

$$R_t = \mu + \sigma_t z_t, \quad z_t \stackrel{\text{i.i.d}}{\sim} (0, 1),$$

so that $\mathbb{E}[R_t] = \mu$ and $\text{var}[R_t] = \sigma_t^2$. In principle, μ could be time-varying as well.

- Let z_p denote the $100p\%$ percentile of


$$z_t = \frac{R_t - \mu}{\sigma_t}.$$

It can be estimated by applying any of the VaR methods above (historical, normal, or Student's t) to the *filtered* (demeaned and devolitized) returns

$$\hat{z}_t = \frac{R_t - \hat{\mu}}{\hat{\sigma}_t}$$

- Finally, $VaR_{t+1}^p = -\mu - \sigma_{t+1} z_p$.

```
In [23]: sig_ewma=np.sqrt(sig2_ewma)
mu=np.mean(r)
z=(r-mu)/sig_ewma
VaR_filtered_hist=-mu-sig_ewma*z.quantile(0.01)
VaR_filtered_hist.plot(color='red');
plt.plot(-r)
plt.title('Filtered Historical VaR');
```



Backtesting

- The Basel accords require that banks' internal VaR models be *backtested*.
- They recommend constructing the 1% VaR over the last 250 trading days and counting the number of *VaR exceptions* (times that losses exceeded the day's VaR figure).
- A method is said to lie in the:
 - Green zone, in case of 0-4 exceptions;
 - Yellow zone, in case of 5-9 exceptions;
 - Red zone, in case of 10 or more exceptions.
- Being in one of the latter two incurs an extra capital charge.

- A more advanced method is the *dynamic quantile* (DQ) test by Engle and Manganelli (2004).
- It is based on the *hit series*

$$I_t = \begin{cases} 1, & \text{if } r_t < -VaR_t^p, \\ 0, & \text{if } r_t > -VaR_t^p. \end{cases}$$

- If the VaR model is correctly specified, then $\mathbb{E}[I_t] = p$ (there should be $p \cdot T$ exceptions in a sample of size T , on average). This is known as the *unconditional coverage hypothesis*.
- It can be tested by regressing $I_t - p$ on an intercept and testing that it is zero.
- In addition, it is desirable that the exceptions not be correlated. This is the *independence hypothesis*. It can be tested by including lags of I_t in the regression and testing their significance.
- Jointly testing both (with an F test) tests the *conditional coverage hypothesis*.

```
In [24]: import statsmodels.formula.api as smf
y=(r<-VaR_filtered_hist)*1 #multiplication by 1 turns True/False into 1/0
y.name='I'
data=pd.DataFrame(y)
model=smf.ols('I.subtract(0.01)~I.shift(1)', data=data)
res=model.fit()
print(res.summary2())
```

Results: Ordinary least squares						
=====						
Model:	OLS		Adj. R-squared:	0.020		
Dependent Variable:	I.subtract(0.01)		AIC:	-2069.9720		
Date:	2017-11-13 20:06		BIC:	-2059.7852		
No. Observations:	1204		Log-Likelihood:	1037.0		
Df Model:	1		F-statistic:	25.67		
Df Residuals:	1202		Prob (F-statistic):	4.68e-07		
R-squared:	0.021		Scale:	0.010475		

	Coef.	Std.Err.	t	P> t	[0.025	0.975]

Intercept	-0.0008	0.0030	-0.2576	0.7967	-0.0066	0.0051
I.shift(1)	0.1446	0.0285	5.0669	0.0000	0.0886	0.2006

Omnibus:	1816.402		Durbin-Watson:	2.014		
Prob(Omnibus):	0.000		Jarque-Bera (JB):	382361.558		
Skew:	9.218		Prob(JB):	0.000		
Kurtosis:	88.334		Condition No.:	10		
=====						

- Conclusions:

- Unconditional coverage is not rejected. This is by construction; note that

$$r_t \lesseqgtr -VaR_t^p \iff z_t \lesseqgtr z_p.$$

- Independence is rejected; apparently our model is dynamically mis-specified.

May need to use a more general GARCH model instead of EWMA.

- The latter finding is likely driving the rejection of the conditional coverage test:

```
In [25]: print(res.f_test('Intercept=0, I.shift(1)=0'))
```

```
<F test: F=array([[ 12.87315967]]), p=2.93962494772e-06, df_denom=1202, df_num=2>
```